

Deep Learning Assignment

Deadline : 11th November 2020

1 Tutorials

Some relevant tutorials to familiarise yourself with pytorch :

- https://pytorch.org/tutorials/beginner/pytorch_with_examples.html
- https://pytorch.org/tutorials/beginner/blitz/cifar10_tutorial.html
- <https://youtu.be/6SlgtELq0Wc?t=3077>
- <https://github.com/kuangliu/pytorch-cifar> (Not a tutorial but repo might help)
- Pytorch Documentation : <https://pytorch.org/docs/stable/index.html>

Jargon you should be comfortable with : Tensors, Optimiser, Models, Autograd or Backward Pass, Activations, ...

2 XOR

2.1 Data Generation

Firstly generate 10K points in $[-1, 1] \times [-1, 1]$ using

```
np.random.seed(0)
data = 2*np.random.uniform(size=(10000,2)) - 1
```

Label points according to their quadrant. First and Third quadrant being labeled 0 and Second and Fourth quadrant being labeled 1

2.2 Exercises

1. Write a Dataset module for the XOR data (3 sets train, validation and test, 70:15:15 respectively)
 - Inherit from `torch.utils.data.Dataset`
 - Define `__init__`, `__getitem__`, `__len__`
2. Define the Dataloader with batchsize of 16
 - Go through all the arguments of dataloader like `drop_last`, `shuffle`, `batch_size`
3. Define the Dataloader with batchsize of 16
4. Define the Neural Network Model
 - Inherit from : `torch.nn.Module`
 - 2 Linear layers, ReLU activation after first layer, single numeric output
 - Variable hidden layer size as input to model (4 as default hidden layer size)
5. Loss function define as `torch.nn.CrossEntropyLoss`
 - Explore other possible loss functions like `LogSoftmax` + `NLLLoss` vs `CrossEntropyLoss` and `MSELoss`
6. Optimizer : Use SGD optimizer with learning rate of 1e-3
 - See `zero_grad()`, `step()`
7. Write the main training loop and validation loops for n epochs
 - See `loss.backward()`, `model.forward()`
 - Use `model.train()`, `model.eval()`, `torch.no_grad()`
 - Validate for each epoch, run for 100 epochs
8. Plot the following :
 - (a) Training and Validation loss vs epoch in a single plot
 - (b) Training and Validation accuracy vs epoch in a single plot
 - (c) Best Validation loss vs Hidden Layer size (use hidden size to be (2,4,6,8,10))
 - (d) Best Validation loss vs learning rate used (use learning rates in (1e-5, 1e-4, 1e-3, 1e-2, 1e-1)) for max number of 20 epochs
 - (e) Plot test set predicted labels for best validation model. Report accuracy and loss for the same.